



UNIVERSIDAD
NACIONAL
DE COLOMBIA

CLASIFICACIÓN Y RECONOCIMIENTO DE PATRONES

Evaluación de Desempeño

John William Branch

Esmeide Alberto Leal Narváez

Agenda



- **Introduction**
- **Classification**
 - Confusion Matrix
 - Precision, Recall, Specificity, F1 Score
 - Precision-Recall AUC
 - ROC/AUC Curve
 - Hold-out, Cross Validation, Leave-one-out (Accuracy Estimation)
- **Examples**
- **Regression**
 - MAE, MSE
 - RMSE, RMSLE, R-Square
- **Examples**

Introduction

- **What is an evaluation metric?**
- A way to quantify performance of a machine learning model
- Evaluation metric \neq Loss function
- There are various metrics which we can use to evaluate the performance of ML algorithms, classification as well as regression algorithms.
- We must carefully choose the metrics for evaluating ML model performance because
 - How the performance of ML algorithms is measured and compared will be dependent entirely on the metric you choose.
 - How you weight the importance of various characteristics in the result will be influenced completely by the metric you choose.

Introduction


- Testing: How well is my model doing?
 - Is good or not
- How do I improve it based on these metrics?



Introduction



- Testing: How well is my model doing?
 - Is good or not
- How do I improve it based on these metrics?

● ● Training
○ ○ Test



Introduction


- Testing: How well is my model doing?
 - Is good or not
- How do I improve it based on these metrics?




● ● Training
○ ○ Test

Introduction

- Testing: How well is my model doing?
 - Is good or not
- How do I improve it based on these metrics?



Generalize



Overfitting - Memorize

● ● Training
○ ○ Test

Confusion Matrix

- A confusion matrix is a predictive analytics tool
- Not a metric.
- Specifically, it is a table that displays and compares actual values with the model's predicted values.
- Within the context of machine learning, a confusion matrix is utilized to analyze how a machine learning classifier performed on a dataset.
- Help to understand some other metrics.
- A confusion matrix generates a visualization of metrics like precision, accuracy, specificity, and recall.

Confusion Matrix

- Formal Definition
 - The confusion matrix, \mathbf{T} , is a $n \times n$ matrix, where n is the number of classes of our data. The element $T(i, j)$ of the confusion matrix is defined as the number of samples that belong to class C_i and were classified as C_j . A perfect classification means that $T(i, i)$ is N_i and $T(i, j) = 0$ for $i \neq j$, where N_i is the number of samples of class C_i .

Confusion Matrix

- Definition

		← Predicted →			
Actual ↑		C_1	C_2	...	C_n
	C_1				
	C_2				
	...				
	C_n				

Confusion Matrix

- Definition

		← Predicted →			
Actual ↑		C_1	C_2	...	C_n
	C_1	1350	80		10
	C_2	77	950		5
	...				
	C_n	30	64		895

Example: there are 80 samples of class 1 that have been classified as class 2.

Confusion Matrix

- Two Class (Detection)

		← Predicted →	
Actual	↑	C_0	C_1
	C_0		
	C_1		

- C_0 = Target Class
- C_1 = Not Target Class
- Example: Disease Detection, Anomaly Detection, Face Detection, etc.

Confusion Matrix

- Toy Example
 - Imagine that we have this diabetes dataset

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

Confusion Matrix

- Toy Example
 - We have some clinical measurements . . .

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1


Confusion Matrix

- Toy Example
 - We want to apply a machine learning method to them to predict whether or not someone will develop diabetes disease.


	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

Confusion Matrix


- To do this, we could use, different machine learning methods




Logistic Regression



Random Forest



SVM



K-NN

etc.

How to decide which one is the best with our data?

Confusion Matrix

- Toy Example
 - We start by dividing the data into Training and Testing sets

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1

Training Data

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
1	1	85	66	29	0	26.6	0.351	31	0





Testing Data

Confusion Matrix

- Toy Example
 - We start by dividing the data into Training and Testing sets

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1

Training Data


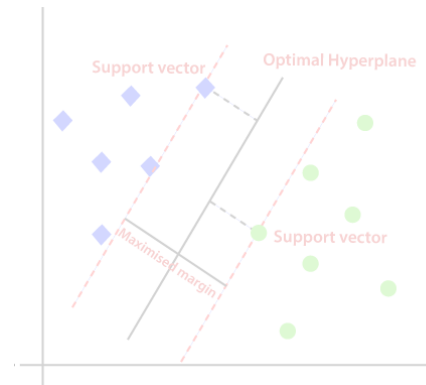




Confusion Matrix

- Toy Example
 - We start by dividing the data into Training and Testing sets

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
1	1	85	66	29	0	26.6	0.351	31	0

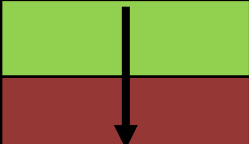
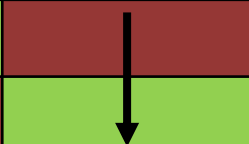
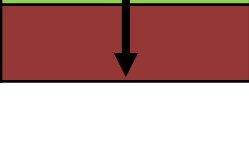
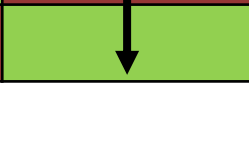
Testing Data



Confusion Matrix

- We create a confusion matrix for each method

← Predicted →

		C_0	C_1
Actual ↑	C_0		
↓	C_1		





- The cols in the confusion matrix correspond to what machine learning algorithm predicted

- C_0 = Has Diabetes
- C_1 = Does Not Have Diabetes

Confusion Matrix

- We create a confusion matrix for each method

← Predicted →

		C_0	C_1
↑ Actual	C_0		
↓	C_1		

- The rows in the confusion matrix correspond to the known truth

- C_0 = Has Diabetes
- C_1 = Does Not Have Diabetes

Confusion Matrix

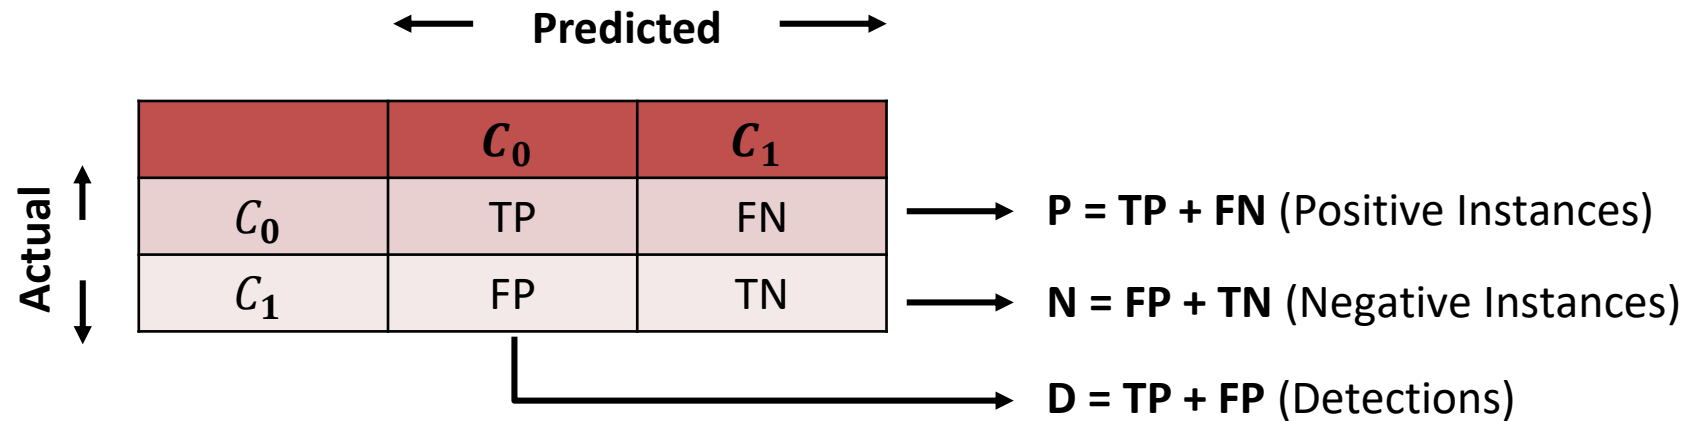
- Two Class

		← Predicted →	
Actual		C_0	C_1
	C_0	TP	FN
	C_1	FP	TN

- **True Positive (TP)**: number of target correctly classified (**Patients with diabetes that were correctly identified by the algorithm**)
- **True Negative (TN)**: number of non-target correctly classified (**Patients without diabetes that were correctly identified by the algorithm**)
- **False Positive (FP)**: number of non-targets classified as targets. The false positives are known as “false alarms” and “Type I Error” (**Patients without diabetes, but the algorithm says they do**)
- **False Negative (FN)**: number of targets classified as non-targets. The false negatives are known as “Type II Error” (**Patients with diabetes, but the algorithm said they don't have diabetes**)

Confusion Matrix


- Two Class



- True Positive (TP):** number of target correctly classified (**Patients with diabetes that were correctly identified by the algorithm**)
- True Negative (TN):** number of non-target correctly classified (**Patients without diabetes that were correctly identified by the algorithm**)
- False Positive (FP):** number of non-targets classified as targets. The false positives are known as “false alarms” and “Type I Error” (**Patients without diabetes, but the algorithm says they do**)
- False Negative (FN):** number of targets classified as non-targets. The false negatives are known as “Type II Error” (**Patients with diabetes, but the algorithm said they don't have diabetes**)

Confusion Matrix

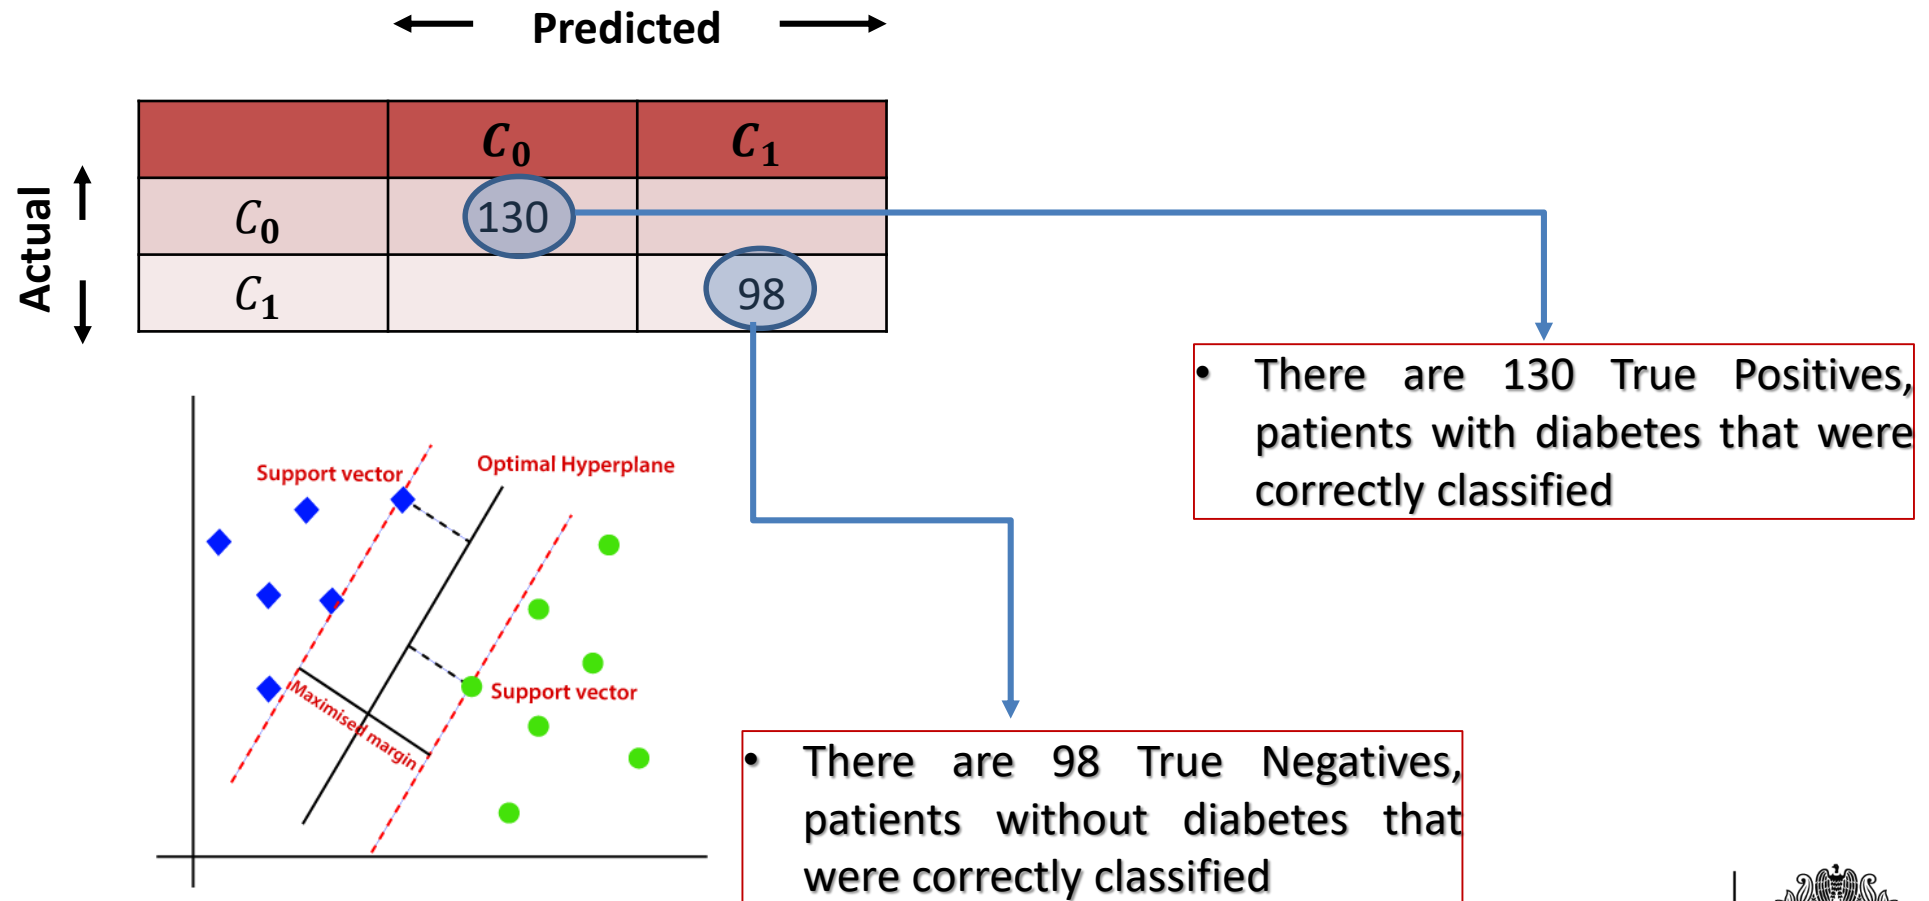
Example



		← Predicted →	
Actual		C_0	C_1
	C_0	TP= 6	FN= 1
	C_1	FP= 2	TN= 5

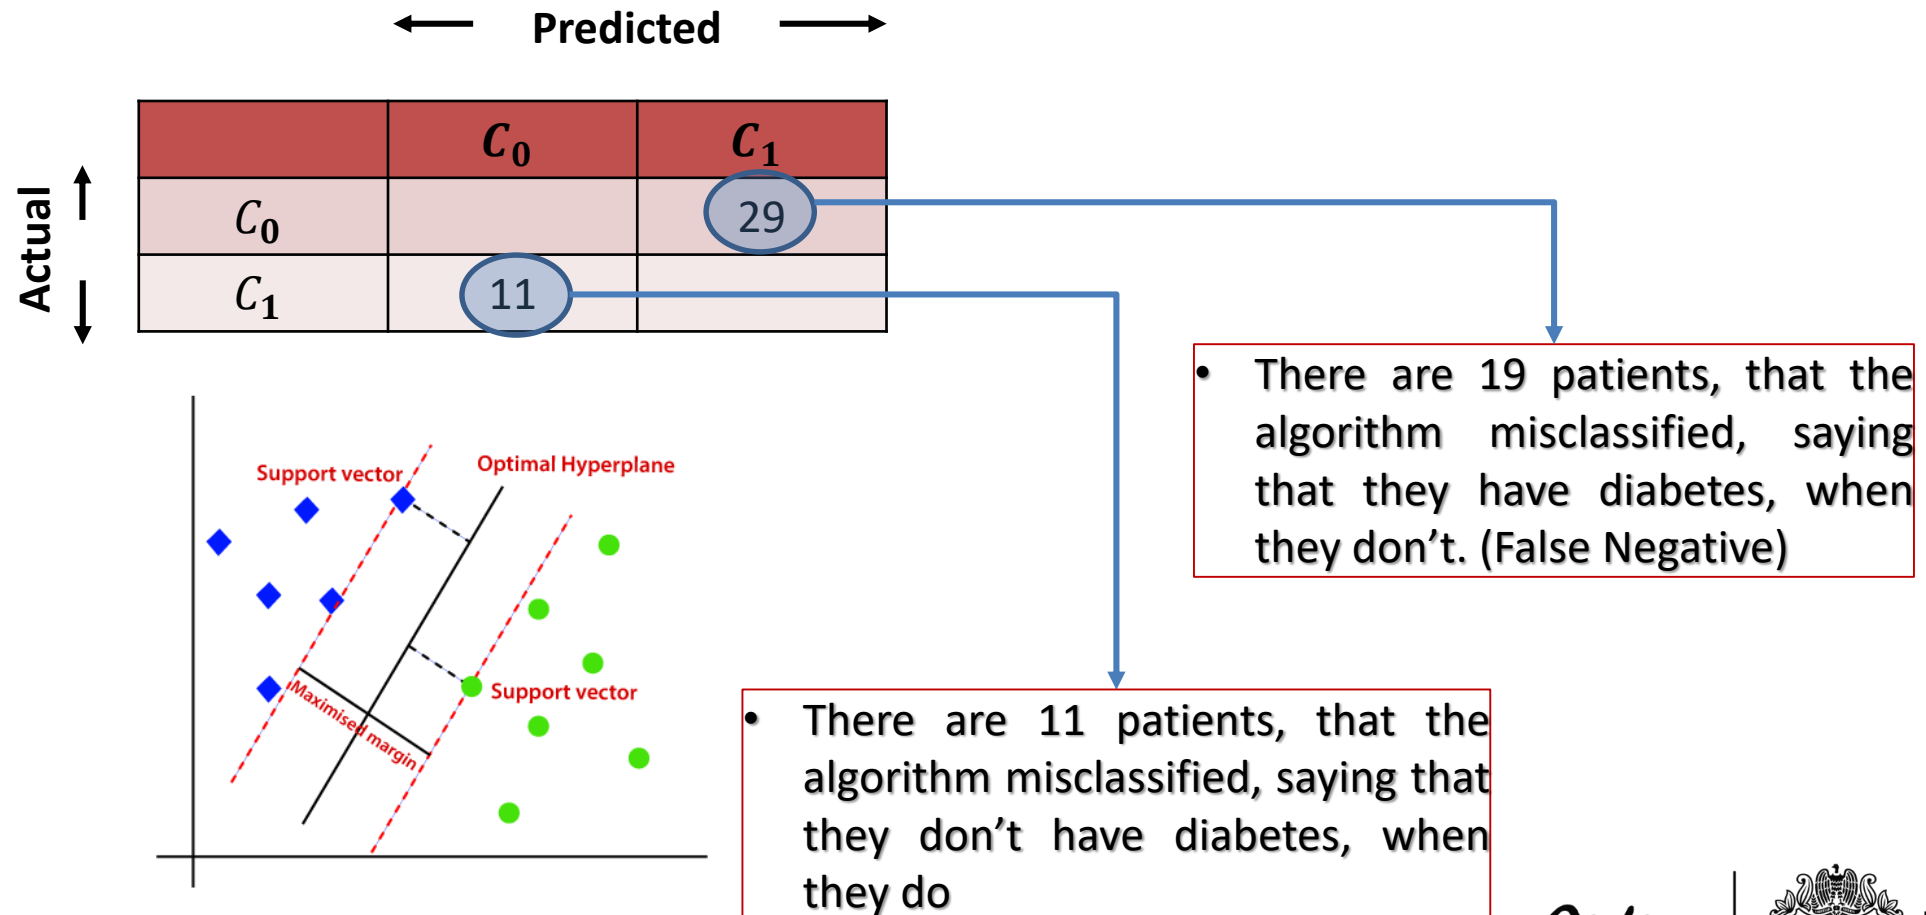
Confusion Matrix

- Example



Confusion Matrix

- Example



Confusion Matrix

- Example


		← Predicted →	
Actual		C_0	C_1
	C_0	130	27
	C_1	14	98

- The diagonal of the confusion matrix (blue cells) indicate how many times the samples were correctly classified.


- The secondary diagonal of the confusion matrix (green cells) are samples that the algorithm misclassified.

Confusion Matrix

- Example




	C_0	C_1
C_0	130	27
C_1	14	98




	C_0	C_1
C_0	106	38
C_1	45	58

Confusion Matrix

- Example



	C_0	C_1
C_0	130	27
C_1	14	98




	C_0	C_1
C_0	106	38
C_1	45	58


- The k-NN algorithm, had worse performance than the SVM algorithm, predicting patients with diabetes.

Confusion Matrix

- Example



	C_0	C_1
C_0	130	27
C_1	14	98




	C_0	C_1
C_0	106	38
C_1	45	58

- The k-NN algorithm, had worse performance than the SVM algorithm, predicting patients without diabetes.

Confusion Matrix

- Example




	C_0	C_1
C_0	130	27
C_1	14	98


- Conclusion: we choose the SVM algorithm, instead The k-NN algorithm.

Confusion Matrix

- Example



	C_0	C_1
C_0	130	27
C_1	14	98




	C_0	C_1
C_0	127	29
C_1	12	101


- Now we use logistic regression to testing the dataset and create its confusion matrix.

Confusion Matrix

- Example



	C_0	C_1
C_0	130	27
C_1	14	98



	C_0	C_1
C_0	127	29
C_1	12	101

- Is difficult to choose what algorithm have best performance, then we use a more sophisticated metrics to help us to take a decision

Confusion Matrix

```
from sklearn.metrics import confusion_matrix

X_actual = [1, 1, 0, 1, 0, 0, 1, 0, 0, 0]
Y_predic = [1, 0, 1, 1, 1, 0, 1, 1, 0, 0]

results = confusion_matrix(X_actual, Y_predic)

print ('Confusion Matrix :')
print(results)
```

Output

```
Confusion Matrix:
[
  [3 3]
  [1 3]
]
```

- Definitions

True positive rate, known as Sensitivity or Recall:

$$TPR = S_n = Re = \frac{TP}{P} = \frac{TP}{TP + FN}$$

Precision or Positive Predictive Value:

$$Pr = \frac{TP}{D} = \frac{TP}{TP + FP}$$

True negative rate, known as Specificity:

$$TNR = Sp = \frac{TN}{N} = \frac{TN}{TN + FP}$$

False negative rate, known as Miss Rate:

$$FNR = MR = \frac{FN}{P} = \frac{FN}{TP + FN}$$


False positive rate, known as 1-Specificity:

$$FPR = 1 - Sp = \frac{FP}{N} = \frac{FP}{TN + FP}$$

Accuracy:

$$ACC = \frac{TP + TN}{P + N}$$

Accuracy



Accuracy:

$$ACC = \frac{TP + TN}{P + N}$$

$$\text{Accuracy} = \frac{\text{Correctly Classified points}}{\text{All points}}$$

$$\text{Accuracy} = \frac{11}{11 + 3}$$

$$\text{Accuracy} = 78.57\%$$

- Accuracy = Out of the all data, how many did we classified correctly?
- Is the correct metric?

Accuracy

- Example
- Imagine the following problem: Suppose we must classify 10.000 peoples and we need to classify as Chinese or not Chinese.
- Now, we have that of these 10.000 people 99,990 are not Chinese and 10 are Chinese
- **Our classifier is a simple program that only returns false when evaluating anyone.**
- Evaluating Accuracy

Accuracy

- From the confusion matrix, we can estimate Accuracy

		← Predicted →				← Predicted →		
Actual		C_0	C_1	Actual		C_0	C_1	
	C_0	TP	FN		C_0	TP	FN	P
	C_1	FP	TN		C_1	FP	TN	N

Accuracy:

$$ACC = \frac{TP + TN}{P + N} = \frac{99,990}{10,000} = 99,9\%$$

- Accuracy = Out of the all the Chinese people, how many did we classified correctly?
- What is wrong with this? and why we need other performance measures?

Accuracy

```
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score

X_actual = [1, 1, 0, 1, 0, 0, 1, 0, 0, 0]
Y_predic = [1, 0, 1, 1, 1, 0, 1, 1, 0, 0]

results = confusion_matrix(X_actual, Y_predic)

print ('Confusion Matrix :')
print(results)
print ('Accuracy Score is',accuracy_score(X_actual, Y_predic))
```




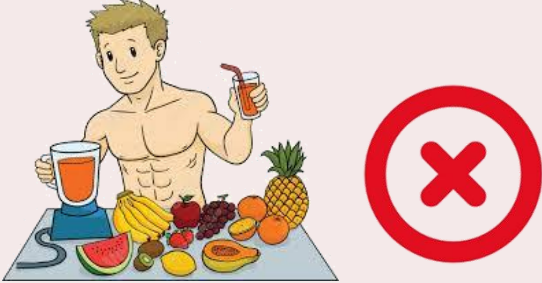
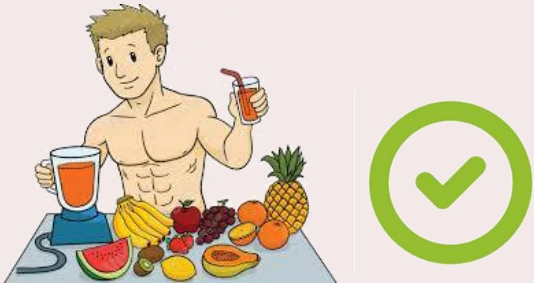
Output

Confusion Matrix:

```
[
  [3 3]
  [1 3]
]
```



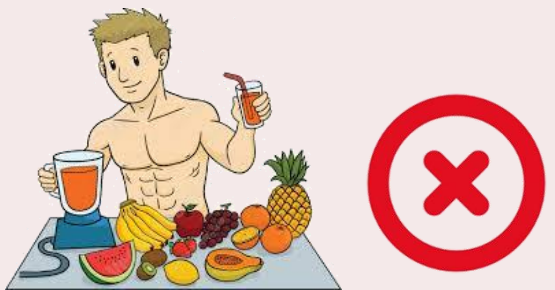
Accuracy Score is 0.6

Precision - Recall

	Diagnosed Sick	Diagnosed Healthy
Sick		
Healthy		






Medical Model

Precision - Recall

	Diagnosed Sick	Diagnosed Healthy
Sick		
Healthy		




Medical Model

Precision - Recall

	Sent to Spam Folder	Sent to Inbox
Spam		
Not Spam		

Spam Detector

Precision - Recall

	Sent to Spam Folder	Sent to Inbox
Spam		
Not Spam		

Spam Detector

Precision - Recall



- **Medical Model**
- False Positives Ok
- False Negative **NOT** Ok
- Find all the sick people
- Ok if not all are sick
- **High Recall**



- **Spam Detector**
- False Positives **NOT** Ok
- False Negative Ok
- You don't necessarily need to find all spam but they better all be spam
- **High Precision**

Precision - Recall

- From the confusion matrix, we can estimate Precision and Recall

		Predicted	
		C_0	C_1
Actual	C_0	TP	FN
	C_1	FP	TN
		D	

Precision or Positive Predictive Value:

$$Pr = \frac{TP}{D} = \frac{TP}{TP + FP}$$

- Precision** = Out of all the people classified Chinese, how many are actually Chinese? (Or classified correctly)

		Predicted	
		C_0	C_1
Actual	C_0	TP	FN
	C_1	FP	TN
		P	

True positive rate, known as Sensitivity or Recall:

$$TPR = S_n = Re = \frac{TP}{P} = \frac{TP}{TP + FN}$$

- Recall** = Out of all the people that are actually Chinese, how many are classified as Chinese?

Precision - Recall

		← Predicted →	
Actual		C_0	C_1
	C_0	TP	FN
	C_1	FP	TN

- Precision = Of the people classified Chinese, how many are actually Chinese?
- Number of people classified Chinese = TP+ FP
- Number of people actually Chinese (when classified as Chinese) = TP
- Recall = Of the people that are actually Chinese, how many are classified as Chinese?
- Number of people actually Chinese = TP+ FN
- Number of people classified Chinese (when actually Chinese) = TP

Precision - Recall

		Predicted	
		C_0	C_1
Actual	C_0	TP = 0	FN = 10
	C_1	FP = 0	TN = 9,990

$$Precision = \frac{TP}{TP + FP}$$

$$Precision = \frac{0}{0} \text{ (not defined)}$$

$$Recall = \frac{TP}{TP + FN}$$

$$Recall = \frac{0}{0 + 10} = 0$$

$$Accuracy = \frac{TP+TN}{TP+TN+FP+FN} = 99,9\%$$

- Now, we redefine the classifier: **Our classifier is a simple program that only returns true when evaluating anyone.**

		Predicted	
		C_0	C_1
Actual	C_0	TP = 10	FN = 0
	C_1	FP = 9,990	TN = 0

$$Precision = \frac{10}{10 + 9,990} = 0,001$$

$$Recall = \frac{10}{10 + 0} = 1$$

$$Accuracy = \frac{TP+TN}{TP+TN+FP+FN} = 0.1\%$$

Precision - Recall

```
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report

X_actual = [1, 1, 0, 1, 0, 0, 1, 0, 0, 0]
Y_predic = [1, 0, 1, 1, 1, 0, 1, 1, 0, 0]

results = confusion_matrix(X_actual, Y_predic)

print ('Confusion Matrix :')
print (results)
print ('Accuracy Score is', accuracy_score(X_actual, Y_predic))
print ('Classification Report : ')
print (classification_report(X_actual, Y_predic))
```

```
from sklearn.metrics import precision_score
from sklearn.metrics import recall_score

X_actual = [1, 1, 0, 1, 0, 0, 1, 0, 0, 0]
Y_predic = [1, 0, 1, 1, 1, 0, 1, 1, 0, 0]

precision = precision_score (X_actual, Y_predic)
recall = recall_score (X_actual, Y_predic)

print ('Precsion Score is', precision)
print ('Recall Score is', recall)
```

Output

Confusion Matrix:

```
[
  [3 3]
  [1 3]
]
```

Accuracy Score is 0.6

Classification Report :

	precision	recall	f1-score	support
0	0.75	0.50	0.60	6
1	0.50	0.75	0.60	4
micro avg	0.60	0.60	0.60	10
macro avg	0.62	0.62	0.60	10
weighted avg	0.65	0.60	0.60	10

F1-Score/-F_β-Score (F-Score/F-Measure)

- Comparing Systems

- System 1
 - Precision: 70%
 - Recall: 60%
- System 2
 - Precision: 80%
 - Recall: 50%

- Which of these two systems perform better?
- The answer to this question is the computing of a single measure from the precision and recall measure, which is called the F beta measure

$$F_{\beta} = \frac{1}{\beta \times \frac{1}{Precision} + (1 - \beta) \times \frac{1}{Recall}}$$

β=0.5 = F-Measure

$$F_{\beta} = \frac{1}{0.5 \times \frac{1}{0.7} + (1 - 0.5) \times \frac{1}{0.6}} = 0.6461$$



$$F_{\beta} = \frac{1}{0.5 \times \frac{1}{0.8} + (1 - 0.5) \times \frac{1}{0.5}} = 0.6153 \quad \text{Medical model}$$

β=0.95

$$F_{\beta} = \frac{1}{0.95 \times \frac{1}{0.7} + (1 - 0.95) \times \frac{1}{0.6}} = 0.6942$$



$$F_{\beta} = \frac{1}{0.95 \times \frac{1}{0.8} + (1 - 0.95) \times \frac{1}{0.5}} = 0.7766 \quad \text{Spam model}$$

Specificity

- From the confusion matrix, we can estimate two metrics: Recall and Specificity

← Predicted →

		C_0	C_1
Actual ↑ ↓	C_0	TP	FN
	C_1	FP	TN

- Specificity tell us what percentage of patients without diabetes were correctly identified
- Of all the people who are healthy, how many of those did we correctly predict?

True negative rate, known as Specificity:

$$TNR = Sp = \frac{TN}{N} = \frac{TN}{TN + FP}$$

Specificity

```
from sklearn.metrics import specificity_score


X_actual = [1, 1, 0, 1, 0, 0, 1, 0, 0, 0]
Y_predic = [1, 0, 1, 1, 1, 0, 1, 1, 0, 0]

specificity = specificity_score (X_actual, Y_predic)

print ('Specificity Score is', specificity)
```

Recall - Specificity

- Example



	C_0	C_1
C_0	127	29
C_1	12	101

- C_0 = Has Diabetes
- C_1 = Does Not Have Diabetes


True positive rate, known as Sensitivity or Recall:

$$TPR = S_n = Re = \frac{TP}{P} = \frac{TP}{TP + FN} = \frac{127}{127 + 28} = 0.81$$

- Recall tell us that 81% of the patients with diabetes were correctly identified

Recall - Specificity

- Example



	C_0	C_1
C_0	127	29
C_1	12	101


True negative rate, known as Specificity:

$$TNR = Sp = \frac{TN}{N} = \frac{TN}{TN + FP} = \frac{101}{101 + 12} = 0.89$$

- Specificity tell us that 89% of the patients without diabetes were correctly identified

Recall - Specificity

- Example



	C_0	C_1
C_0	130	27
C_1	14	98


True positive rate, known as Sensitivity or Recall:

$$TPR = S_n = Re = \frac{TP}{P} = \frac{TP}{TP + FN} = \frac{130}{130 + 27} = 0.83$$

- Recall tell us that 83% of the patients with diabetes were correctly identified

Recall - Specificity

- Example



	C_0	C_1
C_0	130	27
C_1	14	98


True negative rate, known as Specificity:

$$TNR = Sp = \frac{TN}{N} = \frac{TN}{TN + FP} = \frac{98}{98 + 14} = 0.87$$


- Specificity tell us that 87% of the patients without diabetes were correctly identified

Recall - Specificity

- Example



- Recall = 0.81
- Specificity = 0.89



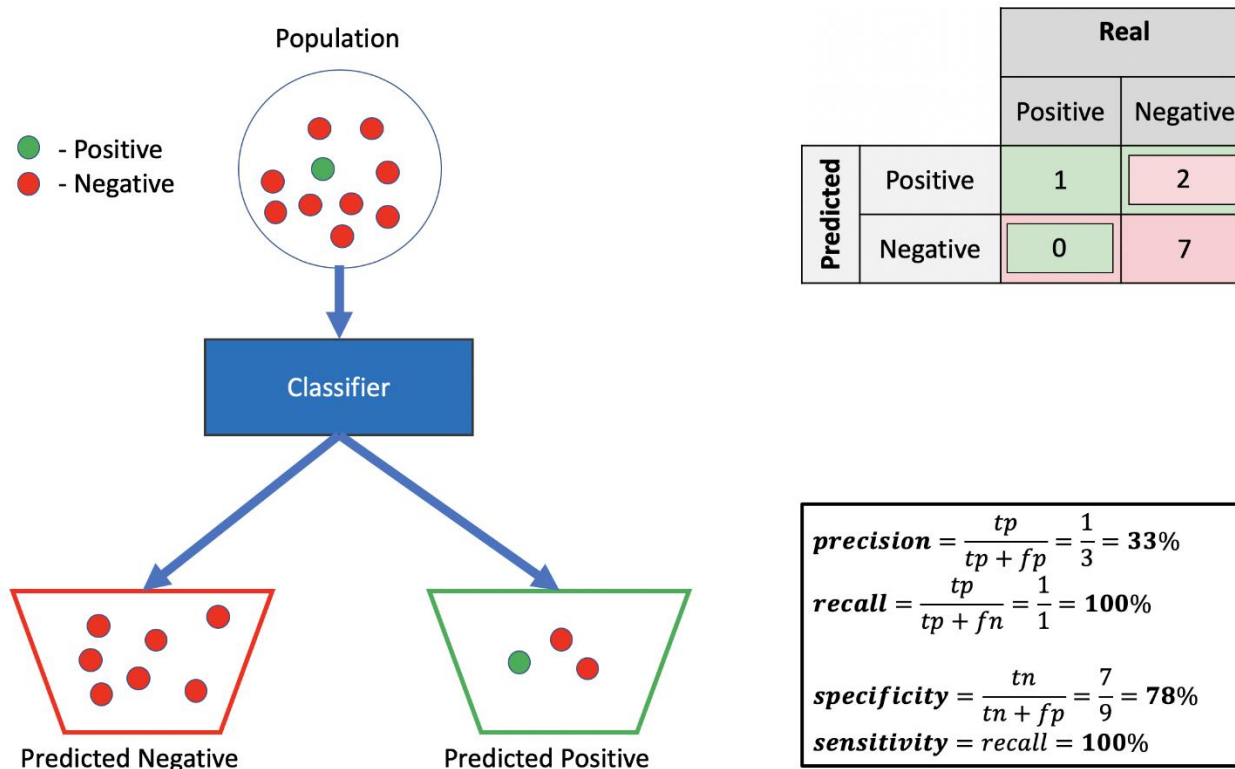
- Recall = 0.83
- Specificity = 0.87

- What to choose?
- We choose the logistic regression model, because correctly identifying patients without diabetes; this is the case if it is more important than identifying patients with diabetes.

- We choose the SVM model, because correctly identifying patients with diabetes; this is the case if it is more important than identifying patients without diabetes.

Recall, Precision, Specificity

Example 1 — Low Precision, High Recall, and High Specificity

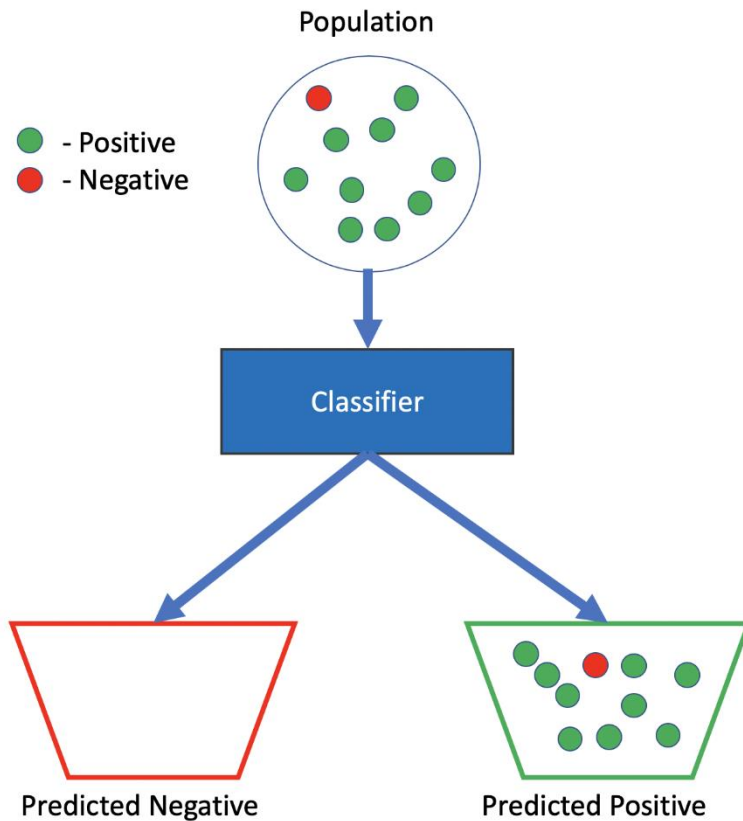


taken from

[Data Science in Medicine — Precision & Recall or Specificity & Sensitivity? | by Alon Lekhtman | Medium](#)

Recall, Precision, Specificity

Example 2 — High Precision, High Recall, and Low Specificity



		Real	
		Positive	Negative
Predicted	Positive	9	1
	Negative	0	0


$$\begin{aligned} \text{precision} &= \frac{tp}{tp + fp} = \frac{9}{10} = \mathbf{90\%} \\ \text{recall} &= \frac{tp}{tp + fn} = \frac{9}{9} = \mathbf{100\%} \\ \text{specificity} &= \frac{tn}{tn + fp} = \frac{0}{1} = \mathbf{0\%} \\ \text{sensitivity} &= \text{recall} = \mathbf{100\%} \end{aligned}$$

taken from

[Data Science in Medicine — Precision & Recall or Specificity & Sensitivity? | by Alon Lekhtman | Medium](#)

Recall, Precision, Specificity

Example 3 — High Precision, Low Recall, and High Specificity



		Real	
		Positive	Negative
Predicted	Positive	1	0
	Negative	8	1

$$\text{precision} = \frac{tp}{tp + fp} = \frac{1}{1} = 100\%$$

$$\text{recall} = \frac{tp}{tp + fn} = \frac{1}{9} = 11\%$$

$$\text{specificity} = \frac{tn}{tn + fp} = \frac{1}{1} = 100\%$$


$$\text{sensitivity} = \text{recall} = 11\%$$

taken from

[Data Science in Medicine — Precision & Recall or Specificity & Sensitivity? | by Alon Lekhtman | Medium](#)

Recall, Precision, Specificity

Example 4 — Low Precision, Low Recall, and High Specificity



		Real	
		Positive	Negative
Predicted	Positive	0	1
	Negative	1	8


$$\begin{aligned} \text{precision} &= \frac{tp}{tp + fp} = \frac{0}{1} = 0\% \\ \text{recall} &= \frac{tp}{tp + fn} = \frac{0}{1} = 0\% \\ \text{specificity} &= \frac{tn}{tn + fp} = \frac{8}{9} = 89\% \\ \text{sensitivity} &= \text{recall} = 0\% \end{aligned}$$

taken from

[Data Science in Medicine — Precision & Recall or Specificity & Sensitivity? | by Alon Lekhtman | Medium](#)

Recall, Precision, Specificity

Example 5 — High Precision, Low Recall, and Low Specificity



		Real	
		Positive	Negative
Predicted	Positive	2	1
	Negative	7	0

$$\text{precision} = \frac{tp}{tp + fp} = \frac{2}{3} = 66\%$$

$$\text{recall} = \frac{tp}{tp + fn} = \frac{2}{9} = 22\%$$

$$\text{specificity} = \frac{tn}{tn + fp} = \frac{0}{1} = 0\%$$

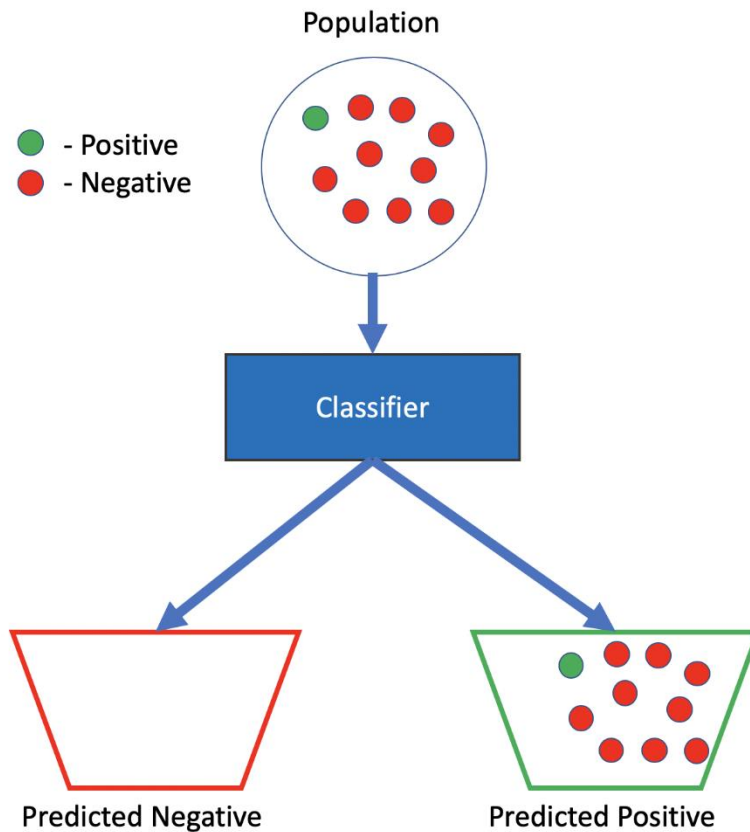
$$\text{sensitivity} = \text{recall} = 22\%$$

taken from

[Data Science in Medicine — Precision & Recall or Specificity & Sensitivity? | by Alon Lekhtman | Medium](#)

Recall, Precision, Specificity

Example 6 — Low Precision, High Recall, and Low Specificity



		Real	
		Positive	Negative
Predicted	Positive	1	9
	Negative	0	0


$$\begin{aligned} \text{precision} &= \frac{tp}{tp + fp} = \frac{1}{10} = 10\% \\ \text{recall} &= \frac{tp}{tp + fn} = \frac{1}{1} = 100\% \\ \text{specificity} &= \frac{tn}{tn + fp} = \frac{0}{9} = 0\% \\ \text{sensitivity} &= \text{recall} = 100\% \end{aligned}$$

taken from

[Data Science in Medicine — Precision & Recall or Specificity & Sensitivity? | by Alon Lekhtman | Medium](#)

Recall, Precision, Specificity

Example 7 — High Precision, High Recall, and High Specificity



		Real	
		Positive	Negative
Predicted	Positive	9	0
	Negative	0	1


$$\begin{aligned} \text{precision} &= \frac{tp}{tp + fp} = \frac{9}{9} = 100\% \\ \text{recall} &= \frac{tp}{tp + fn} = \frac{9}{9} = 100\% \\ \text{specificity} &= \frac{tn}{tn + fp} = \frac{1}{1} = 100\% \\ \text{sensitivity} &= \text{recall} = 100\% \end{aligned}$$

taken from

[Data Science in Medicine — Precision & Recall or Specificity & Sensitivity? | by Alon Lekhtman | Medium](#)

Recall, Precision, Specificity

Example 8 — Low Precision, Low Recall, and Low Specificity



		Real	
		Positive	Negative
Predicted	Positive	0	9
	Negative	1	0

$$\begin{aligned} \text{precision} &= \frac{tp}{tp + fp} = \frac{0}{9} = 0\% \\ \text{recall} &= \frac{tp}{tp + fn} = \frac{0}{1} = 0\% \\ \text{specificity} &= \frac{tn}{tn + fp} = \frac{0}{9} = 0\% \\ \text{sensitivity} &= \text{recall} = 0\% \end{aligned}$$



taken from

[Data Science in Medicine — Precision & Recall or Specificity & Sensitivity? | by Alon Lekhtman | Medium](#)

Summary



- **Accuracy is a great measure** but only **when you have symmetric datasets** (false negatives & false positives counts are close), also, **false negatives & false positives have similar costs**.
- If the cost of false positives and false negatives are different then F1 is your option. **F1 is best if you have an uneven class distribution**.
- **Choose Recall** if the idea of false positives is far better than false negatives.
- **Choose precision** if you want to be more confident of your true positives. for example, Spam emails.
- **Choose Specificity** if you want to cover all true negatives, meaning you don't want any false alarms, you don't want any false positives.

Example



● Class 1



Example





● Class 1 or Target

Blue pixels: Positive instances (P)
White pixels: Negative instances (N)

Example




Example




True positives (TP) : ●
False positives (FP) : ●
False negatives (FN) : ●
True negatives (TN) : ○

Example



IDEAL
TPR = 100%
FPR = 0%

Extreme
No false positive



No false alarm

Reality:
Trade-off between
FPR and TPR






Extreme
All positive samples
are detected



All targets detected

Precision-Recall Curve

- The precision and recall metrics are related so that if you train your classifier to increase accuracy, it will decrease recall and vice versa.
- The PR curve is the result of drawing the graph between precision and recall. This graph allows us to see from which recall we have a precision degradation and vice versa.
- Ideally, a curve that is as close as possible to the upper right corner (high precision and high recall)



True positives (TP) : ●
False positives (FP) : ●
False negatives (FN) : ●
True negatives (TN) : ○



$$Pr = \frac{TP}{D} = \frac{TP}{TP + FP}$$

$$Re = \frac{TP}{P} = \frac{TP}{TP + FN}$$

Precision-Recall Curve

$$Pr = \frac{TP}{D} = \frac{TP}{TP + FP}$$

$$Re = \frac{TP}{P} = \frac{TP}{TP + FN}$$




- AREA UNDER CURVE – AUC:
- Mean Precision Average (Integral)

Precision-Recall Curve

$$Pr = \frac{TP}{D} = \frac{TP}{TP + FP}$$

$$Re = \frac{TP}{P} = \frac{TP}{TP + FN}$$



	Threshold
[[305, 102], [393, 1200]]	0.2
[[263, 144], [210, 1383]]	0.3
[[216, 191], [107, 1486]]	0.4
[[184, 223], [50, 1543]]	0.5
[[109, 298], [10, 1583]]	0.7
[[63, 344], [2, 1591]]	0.8
[[20, 387], [0, 1593]]	0.9

Precision-Recall Curve

```
from sklearn import svm, datasets
from sklearn.model_selection import train_test_split
import numpy as np

iris = datasets.load_iris()
X = iris.data
y = iris.target

# Add noisy features
random_state = np.random.RandomState(0)
n_samples, n_features = X.shape
X = np.c_[X, random_state.randn(n_samples, 200 * n_features)]

# Limit to the two first classes, and split into training and test
X_train, X_test, y_train, y_test = train_test_split(X[y < 2], y[y < 2],
                                                    test_size=.5,
                                                    random_state=random_state)

# Create a simple classifier
classifier = svm.LinearSVC(random_state=random_state)
classifier.fit(X_train, y_train)
y_score = classifier.decision_function(X_test)
```

Compute the average precision score

```
from sklearn.metrics import average_precision_score
average_precision = average_precision_score(y_test, y_score)

print('Average precision-recall score: {0:0.2f}'.format(
    average_precision))
```


Out: Average precision-recall score: 0.88

Precision-Recall Curve

Plot the Precision-Recall curve

```
from sklearn.metrics import precision_recall_curve
from sklearn.metrics import plot_precision_recall_curve
import matplotlib.pyplot as plt




disp = plot_precision_recall_curve(classifier, X_test, y_test)
disp.ax_.set_title('2-class Precision-Recall curve: '
                  'AP={0:0.2f}'.format(average_precision))
```



Out: Text(0.5, 1.0, '2-class Precision-Recall curve: AP=0.88')

ROC Curve – Receiver Operation Characteristic

- The ROC (receiver operating characteristic) curve is similar to the PR curve but changing some values.
- Relate the recall to the false positive rate. In other words, it relates the sensitivity of our model to optimistic failures (classify negatives as positive).
- It makes sense since, generally, if we increase the recall, our model will tend to be more optimistic and will introduce more false positives in the classification.



True positives (TP) : ●
 False positives (FP) : ●
 False negatives (FN) : ●
 True negatives (TN) : ○

$$Pr = \frac{TP}{D} = \frac{TP}{TP+FP}$$



$$Re = \frac{TP}{P} = \frac{TP}{TP+FN}$$

$$FPR = 1 - Sp = \frac{FP}{N} = \frac{FP}{TN+FP}$$

$$TPR = S_n = Re = \frac{TP}{P} = \frac{TP}{TP+FN}$$


- Sensitivity = TPR y Specificity = 1 - FPR

ROC Curve – Receiver Operation Characteristic





- AREA UNDER CURVE – AUC







Precision-Recall Curve and ROC Curve



ROC Curve Example




ROC Curve Example



ROC Curve

```
# roc curve and auc
from sklearn.datasets import make_classification
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_curve
from sklearn.metrics import roc_auc_score
from matplotlib import pyplot
# generate 2 class dataset
X, y = make_classification(n_samples=1000, n_classes=2, random_state=1)
# split into train/test sets
trainX, testX, trainy, testy = train_test_split(X, y, test_size=0.5, random_state=2)
# generate a no skill prediction (majority class)
ns_probs = [0 for _ in range(len(testy))]
# fit a model
model = LogisticRegression(solver='lbfgs')
model.fit(trainX, trainy)
# predict probabilities
lr_probs = model.predict_proba(testX)
# keep probabilities for the positive outcome only
lr_probs = lr_probs[:, 1]
# calculate scores
ns_auc = roc_auc_score(testy, ns_probs)
lr_auc = roc_auc_score(testy, lr_probs)
# summarize scores
print('No Skill: ROC AUC=%.3f' % (ns_auc))
print('Logistic: ROC AUC=%.3f' % (lr_auc))
# calculate roc curves
ns_fpr, ns_tpr, _ = roc_curve(testy, ns_probs)
lr_fpr, lr_tpr, _ = roc_curve(testy, lr_probs)
# plot the roc curve for the model
pyplot.plot(ns_fpr, ns_tpr, linestyle='--', label='No Skill')
pyplot.plot(lr_fpr, lr_tpr, marker='.', label='Logistic')
# axis labels
pyplot.xlabel('False Positive Rate')
pyplot.ylabel('True Positive Rate')
# show the legend
pyplot.legend()
# show the plot
pyplot.show()
```

No Skill: ROC AUC=0.500
Logistic: ROC AUC=0.903

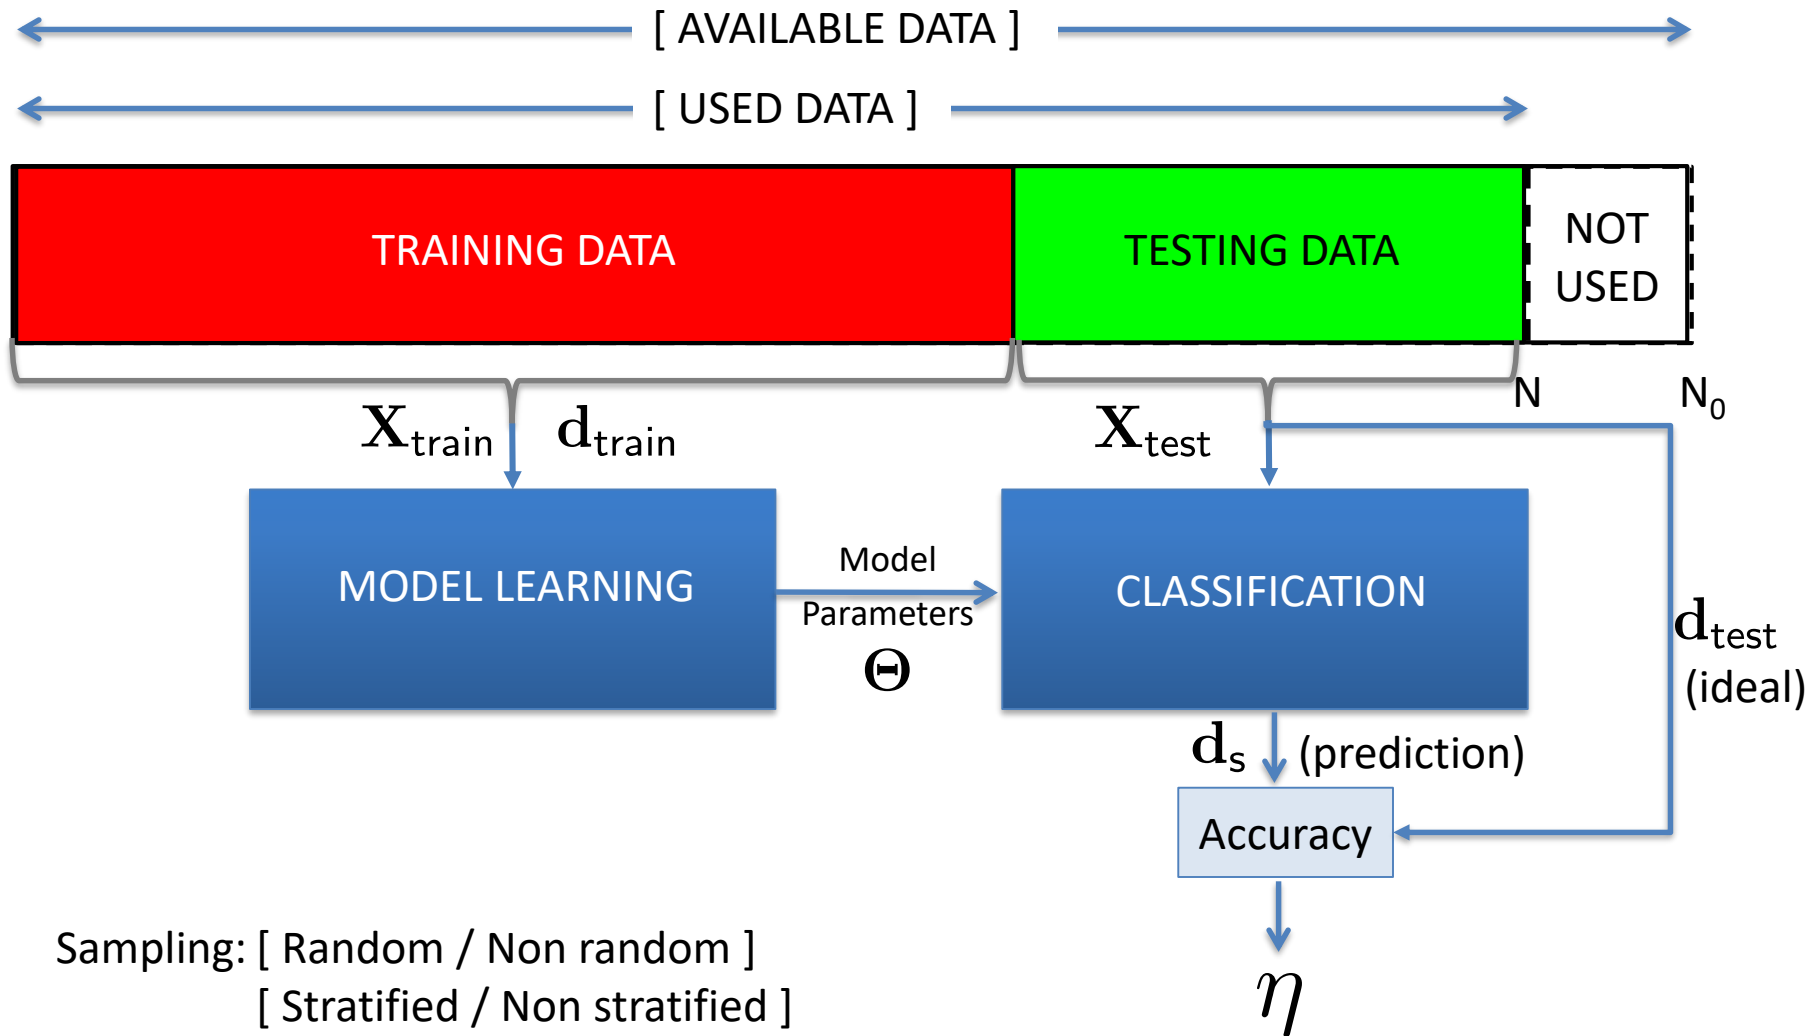


ROC Curve Plot for a No Skill Classifier and a Logistic Regression Model

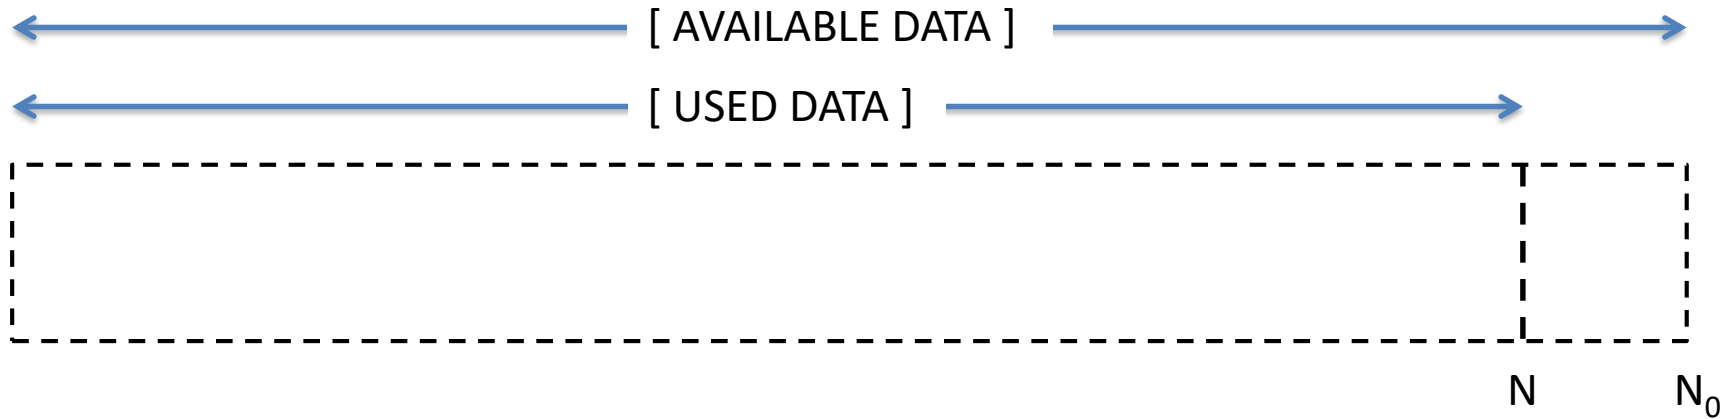
Differences between ROC and PR curves

- In general, we will use the PR curve or the Average Precision when we have problems with unbalanced datasets, that is, when the positive class occurs rarely.
- When there are few positive examples, the ROC curve or the ROC AUC can give a high value, however, the PR curve will be far from its optimal value, revealing an indicator of precision related to the low probability of the positive class.
- It will be an interesting option to use the ROC curve and the AUC ROC when we have a more balanced dataset or we want to reveal an indicator more related to false alarms (false positives).

Estimating the Accuracy of a Classifier



Estimating the Accuracy of a Classifier



Estimating the Accuracy of a Classifier



Estimating the Accuracy of a Classifier




Estimating the Accuracy of a Classifier

HO: HOLD OUT



Estimating the Accuracy of a Classifier

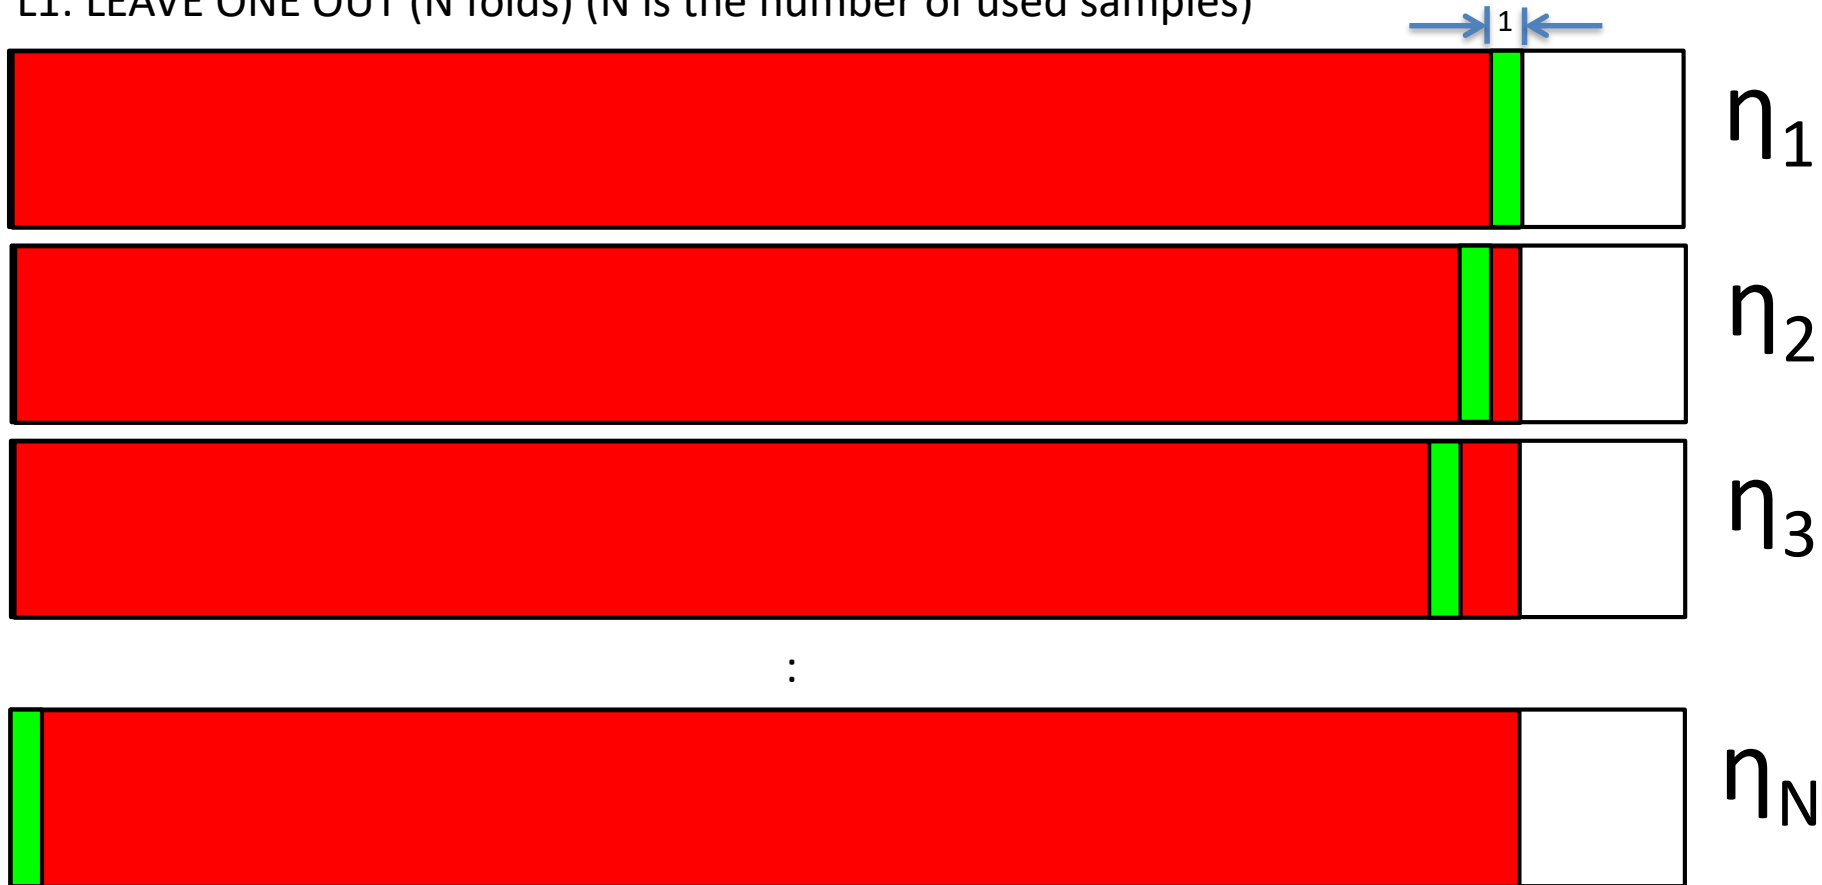
CV: CROSS VALIDATION – n folds



$$\eta = (\eta_1 + \eta_2 + \eta_3 + \dots + \eta_n) / n$$

Estimating the Accuracy of a Classifier

L1: LEAVE ONE OUT (N folds) (N is the number of used samples)



$$\eta = (\eta_1 + \eta_2 + \eta_3 + \dots + \eta_N) / N$$

Estimating the Accuracy of a Classifier


L1*: LEAVE ONE OUT (n folds) (with $n < N$)



$$\eta = (\eta_1 + \eta_2 + \eta_3 + \dots + \eta_n) / n$$

Estimating the Accuracy of a Classifier

```
>> PAT05_Evaluation_HoldOut
All windows closed & all variables deleted.
1)      knn-5    0.9626
2)      knn-7    0.9626
3)      knn-9    0.9626
4)      lda      0.5561
5)      qda      0.9358
6)      svm-lin   0.5989
7)      svm-rbf   0.9572
8)      dmin     0.4439
9)      maha     0.9251
```




Estimating the Accuracy of a Classifier

```
>> PAT05_Evaluation_HoldOut
All windows closed & all variables deleted.
1)      knn-5  0.9626
2)      knn-7  0.9626
3)      knn-9  0.9626
4)      lda    0.5561
5)      qda    0.9358
6)      svm-lin 0.5989
7)      svm-rbf 0.9572
8)      dmin   0.4439
9)      maha   0.9251


>> PAT05_Evaluation_CrossValidation
All windows closed & all variables deleted.
1)      knn-5  0.9467
2)      knn-7  0.9453
3)      knn-9  0.9493
4)      lda    0.6000
5)      qda    0.9053
6)      svm-lin 0.6000
7)      svm-rbf 0.9547
8)      dmin   0.5333
9)      maha   0.9147

>> PAT05_Evaluation_JackKnife
All windows closed & all variables deleted.
Please wait 2-5 minutes...
1)      knn-5  0.9413
2)      lda    0.5987
3)      qda    0.8933
4)      dmin   0.5347
5)      maha   0.8973
```



Regression Metrics

- how to assess model performance?
- After the model fitting, we would like to assess the performance of the model by comparing model predictions to actual (True) data.



$$\text{Residuals}(\text{Error}) = \hat{y}_i - y_i$$

Regression Metrics

- **Mean Absolute Error (MAE):** is obtained by calculating the absolute difference between the model predictions and the true (actual) values.
- **MAE** is a measure of the average magnitude of error generated by the regression model
- **MAE** is calculated as follow:

$$MAE = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

- If **MAE** is zero, this means that the model predictions are perfect

Regression Metrics

- **Mean Square Error (MSE)**: is similar to Mean Absolute Error (**MAE**), but instead of using absolute values, square of the difference between the model predictions and the training dataset (true values) is calculated.
- **MSE** values are generally larger compared to the **MAE** since the residuals are being squared.
- In case of outliers, **MSE** will become much larger compared to **MAE**.
- In MSE, error increases in a quadratic fashion while in **MAE**, is proportional fashion
- **MSE** is calculated as follow:

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

- **Root Mean Square Error (RMSE)**: is the standard deviation of the residuals.
- **RMSE** provide an estimate of how large the residuals are being dispersed
- **RMSE** is calculated as follow:

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

Regression Metrics

- **Mean Absolut Percentage Error (MAPE):** the **MAE** values can range from 0 to infinity which make difficult to interpret the results as compared to the training data.
- **MAPE** is equivalent to **MAE** but provides the error in a percentage form and therefore overcome **MAE** limitations.
- **MAPE** have some limitations if the data point values is zero (there is a division involved)
- **MAPE** is calculated as follow:

$$MAPE = \frac{100\%}{n} \sum_{i=1}^n |y_i - \hat{y}_i| / y_i$$

- **Mean Percentage Error (MPE):** is similar to **MAPE** but without the absolute operation.
- **MPE** is useful to provide an insight equivalent of how many positive errors as compared to negative ones
- **MPE** is calculated as follow:

$$MPE = \frac{100\%}{n} \sum_{i=1}^n (y_i - \hat{y}_i) / y_i$$

Regression Metrics

- **R-Square** or the coefficient of determination represents the proportion of variance (of y) that has been explained by the independent variable in the model.
- If $R^2 = 80$ this means that 80% of the instances in house prices is due to the increase in the size (m^2).

$$R^2 = 1 - \frac{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}{\frac{1}{n} \sum_{i=1}^n (y_i - \bar{y}_i)^2}$$

- In the above equation, numerator is MSE and the denominator is the variance in Y values.

Regression Metrics

```
from sklearn.metrics import r2_score
from sklearn.metrics import mean_absolute_error
from sklearn.metrics import mean_squared_error
X_actual = [5, -1, 2, 10]
Y_predic = [3.5, -0.9, 2, 9.9]
print ('R Squared =', r2_score(X_actual, Y_predic))
print ('MAE =', mean_absolute_error(X_actual, Y_predic))
print ('MSE =', mean_squared_error(X_actual, Y_predic))
```

Output

```
R Squared = 0.9656060606060606
MAE = 0.42499999999999993
MSE = 0.5674999999999999
```